

When Outlier Supervision Fails: Self-Supervised Robustness in Deep Structure-from-Motion

Anonymous CVPR submission

Paper ID *****

Abstract

001 To reconstruct cameras and 3D structure from real im-
002 age collections, learning-based Structure-from-Motion
003 must reject outliers: wrong correspondences that sur-
004 vive geometric verification. The standard recipe trains
005 an outlier classifier with labels derived from an already-
006 trusted reconstruction of the training scenes — su-
007 pervision that presupposes exactly the capability it
008 is meant to provide. We ask when this recipe fails,
009 and show its failures are caused by the labels them-
010 selves, not by domain shift: labels cannot be manufac-
011 tured where robustness is needed most, and retrain-
012 ing the state-of-the-art supervised method (RESfM)
013 on a 60%-contaminated domain with best-case labels
014 still loses $2\times$ to a label-free alternative. That alterna-
015 tive is SS-RESfM: RESfM’s outlier head trained with a
016 self-supervised loss — confident pseudo-labels formed
017 from percentiles of the model’s own reprojection error
018 — with no ground-truth or COLMAP labels anywhere
019 in training or test. Under a strictly controlled protocol
020 (identical tracks, robust bundle adjustment, and train-
021 ing schedule, varying only the outlier mechanism) we
022 compare label-free statistical removal (madweight) and
023 soft reweighting against a from-scratch RESfM base-
024 line across six datasets spanning 0.5%–61% measured
025 outliers. The optimal mechanism flips with contami-
026 nation. Removal wins the high-contamination regime:
027 at 61% plain removal beats the supervised baseline on
028 every seed pair (12.4 ± 2.1 vs. 19.2 ± 3.4) and both
029 beat COLMAP, while at 43% it beats every super-
030 vised seed’s median (1.4 – 1.9 vs. 2.4 – 4.3). Reweighting
031 wins the low-contamination regime, beating the super-
032 vised baseline 2.5 – $2.9\times$ on two of the three clean out-of-
033 distribution splits (BlendedMVS and Olsson; Strecha
034 goes to supervision on every seed) — and, unlike any
035 supervised recipe we test, does so reliably: the label-
036 free arms land in a narrow band on every training
037 seed, whereas the supervised baseline collapses on the
038 clean benchmark in most seeds under every learn-

ing rate, budget, and recipe variant we try (on Ols-
son — 39 essentially outlier-free scenes, 0.5% — ev-
ery MegaDepth-trained supervised variant lands at a
translation error of 7.4–9.1, versus 2.9–3.6 for every
label-free arm). Supervision retains the in-distribution
regime and the cleanest OOD mean (self-supervised
arms within $\sim 38\%$ in distribution). Every compari-
son is backed by five-training-seed bands, a paired-seed
evaluation control, and an environment-noise study —
in a benchmark family where prior work fixes a single
seed. We trace this to a structural circularity: out-
lier labels require a trusted reconstruction of the very
scenes whose contamination makes reconstruction hard
(for any gold-standard labeler, not just COLMAP),
and we measure the resulting label ceiling: regenerat-
ing supervision from scratch, label recall collapses from
0.82 to ~ 0.3 above $\sim 15\%$ contamination, precisely
where robustness is needed. Label-free adaptivity is
thus a missing ingredient for robust deep SfM.

1. Introduction

Structure-from-Motion (SfM) recovers camera poses
and 3D structure from unordered image collections.
Its correspondences are heavily contaminated by out-
liers (wrong matches that survive pairwise geometric
verification), particularly in Internet photo collections
where repeated and symmetric structure produces spu-
rious epipolar geometries. Classical pipelines reject
outliers with RANSAC and robust bundle adjustment
(BA); recent deep equivariant methods (ESfM [2],
RESfM [1]) instead learn a sets-of-sets permutation-
equivariant network that regresses cameras from point
tracks and, in RESfM, adds a supervised per-point in-
lier/outlier classification module. RESfM’s labels are
derived from a COLMAP [4] reconstruction in two
stages: first by track membership, then refined by a
4-pixel reprojection test against COLMAP-pose trian-
gulation [1].

076 We take a different route. Rather than supervise
077 the head with COLMAP labels, SS-RESfM trains it
078 self-supervised: an adaptive confidence-weighted loss
079 forms confident pseudo-labels from percentiles of the
080 model’s own reprojection error, so no ground-truth or
081 COLMAP outlier labels are used anywhere in train-
082 ing or test. On top of this self-supervised head we
083 study a question RESfM’s single supervised mech-
084 anism leaves open: is one outlier-handling mechanism
085 optimal across the range of contamination levels SfM
086 encounters? We find that it is not. Holding every-
087 thing else fixed and varying only the outlier-handling
088 mechanism, the best mechanism flips with contam-
089 ination: hard removal at high contamination, soft
090 reweighting at low, with the supervised classifier re-
091 taining the typical in-distribution scene and the clean-
092 est OOD mean (Fig. 1). The removal family wins
093 the high-contamination regime outright: plain removal
094 beats the supervised baseline on every seed at both
095 43% and $\sim 60\%$, while madweight — MAD-based re-
096 moval with soft-weighted survivors, no learned thresh-
097 old — matches the baseline there and beats COLMAP.
098 The reason is structural, not incidental: outlier labels
099 must be derived from a trusted reconstruction of the
100 very scenes whose contamination makes reconstruction
101 hard, so supervision bootstraps from precisely the ca-
102 pability it is meant to provide — and this circularity
103 binds any gold-standard labeler, not just COLMAP.
104 A self-supervised head consumes only the point tracks;
105 training data for robust SfM becomes as unlimited as
106 unlabeled imagery, and new domains need no labeling
107 infrastructure.

108 Contributions. (1) The label ceiling, measured and
109 shown causal: outlier supervision requires a trusted re-
110 construction of the very scenes whose contamination
111 makes reconstruction hard. We measure the result-
112 ing label collapse (recall $0.82 \rightarrow \sim 0.3$ above $\sim 15\%$ con-
113 tamination) and prove it is causal, not domain shift:
114 retrained on a 60%-contaminated domain with best-
115 case labels, supervision still loses $2\times$ to label-free re-
116 moval (§5.5). (2) The complementarity finding: under
117 an identical tracks+BA+schedule pipeline isolating the
118 mechanism alone, the optimal outlier mechanism flips
119 with contamination — removal at high contamination
120 (beating the supervised baseline on every seed pair at
121 $\sim 60\%$), soft reweighting on clean OOD data (2.5--
122 $2.9\times$, seed-reliably) — with supervision retaining a cu-
123 rated middle band. (3) SS-RESfM, the instrument of
124 the study and a constructive answer: RESfM’s outlier
125 handling made entirely self-supervised via confident
126 pseudo-labels from the model’s own reprojection-error
127 percentiles — no ground-truth or COLMAP labels any-
128 where — matching the supervised baseline where la-

129 bels are sound and surpassing it where they are not;
130 plus madweight, which needs not even the head, remov-
131 ing by robust reprojection statistics (MAD) and beat-
132 ing COLMAP where COLMAP’s own labels cannot be
133 manufactured. (4) A seed- and environment-level reli-
134 ability methodology for post-BA SfM benchmarks: five-
135 training-seed bands for every cell, a paired-seed eval-
136 uation control (2,352 evaluations showing the evalua-
137 tion seed is immaterial), an environment-noise floor for
138 post-BA means, and per-seed-pairing claim verification
139 — practices absent from this benchmark family, where
140 prior work fixes a single seed, and which surface find-
141 ings (bimodal clean-OOD failures, basin-dependent es-
142 capes) that $\text{mean} \pm \text{std}$ alone would hide. (5) A reusable
143 robustness benchmark: rebuilt Appendix-C-protocol
144 track suites for five OOD datasets with measured con-
145 tamination (0.5–61%), GT-derived outlier labels, per-
146 domain train/test splits, and the label-quality bench-
147 mark — released with code and configurations.

2. Related Work 148

149 Deep equivariant SfM. Learning-based SfM re-
150 gresses cameras from point tracks with permutation-
151 equivariant networks. ESfM [2] casts uncalibrated
152 SfM as deep projective factorization over a measure-
153 ment tensor, trained with an unsupervised reprojec-
154 tion loss; GASfM [3] uses a graph-attention variant.
155 Both assume outliers were already removed by two-
156 view geometric verification. RESfM [1] targets contam-
157 inated internet collections with a sets-of-sets architec-
158 ture plus a supervised per-point inlier/outlier classifica-
159 tion module and a robust BA, reporting accuracy close
160 to its own clean-track upper bound (“ESfM*”) and on
161 par with classical Theia/GLOMAP (often with better
162 translation recovery, while registering fewer images).
163 End-to-end pointmap and pose-regression pipelines
164 (VGGSfM, MAST3R, DUST3R) target far smaller im-
165 age sets. Concurrently, VGPA [10] moved robust-
166 ness to the pose-averaging stage with a self-supervised
167 relative-pose consistency objective — independently
168 corroborating that label-free robustness is the frontier
169 of deep SfM, but at a different stage and input modal-
170 ity (pairwise relative poses vs. our point-track matrices
171 with per-observation outliers), so direct comparison is
172 confounded. We propose no new architecture: we reuse
173 the RESfM/ESfM backbone and isolate the outlier-
174 handling stage and its supervision, with a reliability
175 methodology that applies to either paradigm.
176 Outlier rejection, robust BA, and TTT. Correspond-
177 ence outliers that survive RANSAC/two-view filtering
178 are the central difficulty in internet-scale SfM; classical
179 pipelines suppress them with M-estimators and robust
180 BA, RESfM casts rejection as supervised classifica-

tion. We contrast that classification against label-free statistical removal via the median absolute deviation (MAD) [5] and soft reweighting, and show their relative merit is contamination-dependent. RESfM’s robust BA (outlier-point removal, largest connected component, re-triangulation, re-optimization) is applied identically across all arms, and per-scene test-time fine-tuning is treated as an orthogonal option ($\pm$ TTT) applied uniformly.

3. Method

3.1. Background

The input is a (sparse) point-track tensor M of m tracks across n cameras, built by concatenating pairwise RANSAC-verified SIFT matches. The sets-of-sets permutation-equivariant network [1, 2] maps M to per-camera poses and a per-point outlier score $s \in [0, 1]$ (points in the same track may receive different scores). Crucially, unlike RESfM, whose classification head is trained with a BCE loss against COLMAP-derived inlier/outlier labels, SS-RESfM trains this head self-supervised, with an adaptive confidence-weighted loss: per forward pass it forms confident pseudo-labels from percentiles of the model’s own reprojection error (lowest $p\%$ inliers, highest $p\%$ outliers) and applies BCE only to those. No ground-truth or COLMAP outlier labels are used at any point. A reprojection error e is available for every observation once an initial reconstruction exists. All mechanisms below consume s (and, for madweight, e), produce a per-observation weight w (or a hard keep/drop decision), and feed the result to the identical robust-BA stage (Fig. 2).

3.2. Outlier-handling mechanisms

We compare four mechanisms, illustrated as weight-vs-signal transfer functions in Fig. 2:

- remove (the RESfM analogue): drop a point if $s > \tau$; keep with $w=1$ otherwise. We use RESfM’s default $\tau=0.6$ and its $\tau=0.8$ for the low-outlier, single-camera Strecha/BlendedMVS sets [1].
- weight: soft, $w = 1 - s$, no removal.
- hybrid: three-band: keep (s below a low percentile), soft-weight (middle), drop (s above a high percentile).
- madweight (ours): removal is driven by robust reprojection statistics, not the head. An observation is removed if its reprojection error exceeds $\text{med}(e) + \alpha \cdot \text{MAD}(e)$ ($\alpha=2$); surviving observations are soft-weighted by $w = 1 - s$. madweight is the only mechanism that consumes the reprojection branch e and requires no learned threshold for removal.

Table 1. Datasets: outlier rates (RESfM-reported vs. measured on our tracks) and regime.

Dataset	RESfM% ours%		Regime
Strecha [8]	2.7	1.7	clean, LIDAR GT
BlendedMVS [9]	3.6	3.1	low-outlier
MegaDepth [7]	25.6	25.4	contam., in-dist.
1DSfM [6]	31.9	43.3	contam., OOD
1DSfM-hard	–	~ 60	extreme, OOD
Olsson [11]	–	0.5	outlier-free

Each mechanism is evaluated with and without test-time fine-tuning ($\pm$ TTT), in which the head continues to adapt on the test scene.

3.3. Test-time fine-tuning and robust BA

We load a multi-scene checkpoint, run a per-scene forward pass to obtain s and e , apply the mechanism, fine-tune per scene (~ 1000 iterations) under the reprojection objective, and apply robust BA. This protocol, the seed, and the schedule are held fixed across all arms; only the mechanism block changes.

4. Experiments

4.1. Datasets and contamination

We evaluate on six datasets spanning a wide contamination range. Alongside the RESfM paper’s quoted outlier fractions we report the contamination measured on our own tracks (robust ground-truth-pose triangulation with the standard 4px cutoff, applied uniformly to every dataset); see Table 1. On MegaDepth (tracks identical to RESfM’s) the two measurements agree (25.4 vs 25.6), validating the measurement; our 1DSfM rebuild (43.3%) is markedly harder than the tracks behind RESfM’s quoted 31.9%. This yields a falsifiable prediction: since the winning mechanism is a function of contamination (§5.5), on the authors’ own $\sim 32\%$ 1DSfM tracks we expect supervision to be competitive or ahead, with the label-free advantage emerging toward 43% and dominating by 60%. We state this in advance of evaluating on those tracks. 1DSfM-hard = the five scenes unused by the standard ten-scene split, same pipeline, 300 cameras per scene. Networks train on 27 MegaDepth scenes; MegaDepth test is in-distribution, all other datasets are OOD.

4.2. Protocol, baselines, and track provenance

Following RESfM’s public code and default configuration, all networks train for 20,000 epochs (batch size, image subsampling, the eleven-milestone decay schedule, and validation-based early stopping identical across arms), with five training seeds (20–24) per

cell. Two deliberate settings are disclosed. First, the base learning rate is per-method: 10^{-3} for the supervised arms (RESfM’s released value), 10^{-4} for the self-supervised arms, whose pseudo-label loss converges early (best checkpoint $\sim 7.5\text{k}$ vs. $\sim 16.5\text{k}$ epochs); neither is truncation-limited. Second, the infinity-point margin and the checkpoint-selection metric are held identical across arms rather than at RESfM’s released values — a deviation that demonstrably favors the baseline: under the authors’ own Accuracy-based selection it is weaker on every dataset (RESfM-faithful row in Table 2; MegaDepth 0.50 vs. 0.34), so all supervised-baseline margins are lower bounds, and its clean-benchmark failures persist under its own rule (Olsson 8.8, BlendedMVS rotation 53°). We report shallow (1×3) and deep (2×3) architectures.

Fair-comparison anchor. Our controlled baseline is RESfM-from-scratch: trained by us with the identical setup as our arms (same 27 scenes, sampling, epochs, seed), so only the outlier mechanism differs. We additionally report RESfM-released (authors’ weights) and RESfM (paper).¹

Track provenance. MegaDepth tracks and the 36 test scenes are identical to RESfM (directly comparable to published numbers). The OOD tracks are built by us from the released images and ground truth (RESfM did not release OOD tracks), so OOD comparisons are controlled within our pipeline and not directly comparable to RESfM’s published OOD numbers. RESfM’s released weights yield 0.209 on MegaDepth in our harness (vs. reported 0.169; unchanged, 0.217, under the authors’ exact pinned software stack), isolating the residual gap to solver-level nondeterminism in robust BA (supplementary); the stack-replication study there shows the no-TTT arms are insensitive to environment, while TTT means inherit solver-level noise.

A method-isolated comparison, not a SOTA claim. RESfM-from-scratch shares every pipeline component with our arms and differs only in the outlier mechanism — the same controlled protocol RESfM uses when re-training ESfM/GASfM in its pipeline. The anchor is also authors-faithful: five retrains under the authors’ exact validation metric reproduce its five-seed band on every dataset (supplementary), so the comparison does not depend on which supervised recipe one prefers. We claim no SOTA: the released checkpoint (0.209 on MegaDepth) beats either from-scratch arm, but that gap stems from the released weights themselves (the authors’ own training code on our

¹The ESfM row likewise follows the unified 20k-epoch protocol. ESfM’s own released multi-scene (calibrated) recipe trains for 30k epochs plus a 500-epoch per-scene fine-tune; we deliberately hold the schedule fixed across all rows so that the outlier mechanism remains the only controlled variable.

tracks reproduces 0.34–0.39), handicapping supervised and self-supervised arms equally; RESfM-released and RESfM (paper) are reference rows throughout.

Metrics. Following RESfM we draw conclusions from post-BA mean errors (translation emphasized; rotation in degrees; translation is the aligned camera-position error in scene coordinates), with medians as supporting statistics — except on 1DSfM, whose mean is dominated by two symmetric-structure failure scenes (Ellis Island, Tower of London) that fail for all methods; there the median is representative. Rotation and registered-camera counts (N_r) for all arms are tabulated in the supplementary; rotation broadly agrees with translation (the supervised BlendedMVS failure is even larger in rotation), and under high contamination the label-free arms register more cameras than the supervised baseline (38–43% vs. 37% on 1DSfM-hard). Two registration notions must be distinguished: every learned arm outputs a pose for every camera (pose coverage; COLMAP genuinely fails on 36% of 1DSfM-hard), while N_r counts the cameras surviving robust BA and a 5-pixel quality mask, the set over which errors are computed. Under this stricter notion the plain-removal arms attain their low errors over a reduced set (59%/28% vs. the baseline’s 77%/37% on 1DSfM/1DSfM-hard), consistent with aggressive pruning disconnecting weakly-supported cameras; madweight is registration-matched to the baseline (74%/38%), so its comparisons are not coverage artifacts.

5. Results

5.1. Main comparison

Table 2 is the controlled head-to-head under the RESfM-aligned schedule: every row shares tracks, architecture, schedule, robust BA, and seed, so any difference is attributable to the outlier mechanism and its supervision alone. Reading along the contamination axis: (1) Clean OOD (Olsson 0.5%, Strecha 1.7%, BlendedMVS 3.1%): the label-free arms win seed-robustly where the supervised head misfires: every arm beats the full baseline band on BlendedMVS (0.14–0.24 vs. 0.35 ± 0.04 ; soft arms $2.5\times$ better, and the baseline’s 31.7° rotation failure persists on every seed), and by $\sim 3\times$ on Olsson against four of five baseline seeds (3.0 ± 0.3 vs. 8.2–8.9; one seed escapes at 2.69 — the collapse is basin-dependent, frequency 4/5). Figure 3 shows what it looks like: on an essentially outlier-free scene the supervised reconstruction degenerates to a compressed cluster (16.5° rotation), while the label-free arm reproduces the ground-truth arc nearly exactly. Nor is the collapse a from-scratch artifact: the

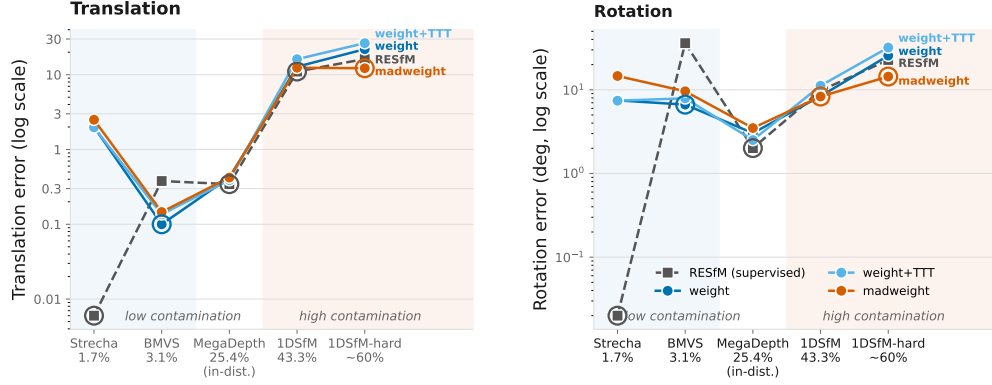


Figure 1. The optimal outlier mechanism flips with contamination. Translation (left) and rotation (right) means (log scale) of the self-supervised arms vs. supervised RESfM, five of the six datasets ordered by measured contamination (Table 2 data). Circled = best per dataset. The rotation panel mirrors the translation ordering, with a self-supervised arm best on 1DSfM and the supervised BlendedMVS failure even more pronounced (36°).




Figure 2. Method. A shared pipeline (top) in which only the outlier-mechanism block changes. The four mechanisms as weight-vs-signal transfer functions (bottom). Only madweight consumes the reprojection-error branch e , and only for its removal step; its surviving points are still weighted by the head score s ($w = 1 - s$). The other three mechanisms use s alone.

released checkpoint also collapses on Olsson (9.10; 9.18 under the authors' pinned stack), as does the authors-faithful configuration (8.77). Strecha is the exception: RESfM's mean wins on every seed (0.006–1.67 vs. ~ 2), though the best label-free median ties it (0.005) at seed 20. (2) In distribution (MegaDepth), supervised RESfM retains mean and median. Its mean is strongly seed-dependent (0.37 ± 0.12 over five seeds; its best, 0.19, approaches the released checkpoint) while its median is stable (0.07 ± 0.02); the deterministic self-supervised arms band at 0.51 ± 0.08 and the TTT arm at 0.40 ± 0.11 over five seeds; a $\sim 38\%$ band-level gap

without labels, not a win (against the authors-faithful configuration, 0.496, the arms reach parity, and the TTT arm is ahead). Medians tell the same story more tightly (0.13–0.20 vs. 0.07). (3) Moderate contamination (1DSfM, 43%) has no stable winner between madweight (11.0 ± 1.7) and the baseline (10.5 ± 1.5 ; overlapping bands), so we decline a directional claim for madweight. The removal family leads on both statistics in band terms: remove+TTT 8.9 ± 1.2 vs. the baseline's 10.5 ± 1.5 (separated means, overlapping tails), and plain remove's medians (1.7 ± 0.2) beat every baseline seed pairwise at five seeds each (1.4–

1.9 vs. 2.4–4.3; its means, 9.1 ± 2.2 , remain failure-scene-dominated): a label-free win at 43% (the supplementary matrix table), established against the baseline’s own five-seed band. (4) Extreme contamination (1DSfM-hard, $\sim 60\%$): removal beats weighting robustly, madweight (18.9 ± 4.1) statistically matches the baseline (19.2 ± 3.4), and plain remove goes further: 12.4 ± 2.1 , below the baseline on every one of the 25 seed pairs (worst seed 15.8 vs. the baseline’s best 16.2; medians 10.8 ± 2.8 vs. 16.5 ± 4.8), and below the released checkpoint (15.29) — the extreme-contamination win holds against every supervised variant we can construct or obtain, and all label-free arms are far ahead of COLMAP (§5.4). Rotation (Fig. 1, right) follows the same ordering.

Classical baselines. Running GLOMAP (global) and COLMAP (incremental) on the same tracks with known intrinsics reproduces the classical/deep trade-off this literature reports, and sharpens it. GLOMAP is far the best method on pristine calibrated data (Strecha 0.047, an order of magnitude below every learned method; Olsson median 0.087) and registers almost everything (96–100% of cameras). Under contamination its translations collapse (3.3 on MegaDepth, 27.2 at 43%, 35.4 at 60%, vs. 9.1 and 12.4 for label-free removal) while its rotations stay strong (8.9° at 43%, better than the supervised baseline’s 9.6°). The failure of classical global SfM under outlier contamination is therefore specifically a positioning failure, not a loss of orientation structure, which is also why translation is the metric of record in this setting. The comparison is not compute- or coverage-matched: GLOMAP keeps 96% of cameras at 43% contamination where plain removal keeps 58%, so the learned arms are far more accurate on a smaller retained set.

5.2. The complementarity finding

Fig. 1 shows Table 2 as a picture, and one pattern organizes it. Between supervision regimes, the self-supervised wins concentrate exactly where a label-free design should help: on clean OOD data, where a classifier trained on contaminated MegaDepth misfires (on BlendedMVS, RESfM’s rotation error explodes to 36°), and at extreme contamination, where supervision itself can no longer be manufactured (§5.3). Within the self-supervised family, the removal-vs-weighting ordering is monotone in contamination: weighting is best at 2–3%, the two tie in distribution, and removal dominates from 43% upward (madweight 12.4 vs. weight 12.8, plain remove 8.9, the supplementary matrix table; 12.4 ± 2.1 vs. ~ 23 at 60%). This mechanism-level flip, the answer to our central question, is stable across schedules and seeds.

Baseline learning-rate robustness, the full mechanism matrix, the MAD/STD/Huber ablation, per-scene selection, and qualitative results are in the supplementary.

5.3. Why label-free removal wins under high contamination

Two factors explain the erosion of the supervised advantage. First, domain shift: the classifier’s decision boundary, fit to MegaDepth, does not transfer to OOD outlier statistics — and even in distribution its classification is mediocre by RESfM’s own ablation (outlier recall 72.1% at precision 44.1%, their Table 7), with their Appendix D crediting much of the final accuracy to robust BA, consistent with our pre/post-BA analysis (supplementary): much of the supervised system’s robustness is BA repair, not classification. Second, and more fundamentally, the supervision itself degrades: RESfM’s labeling assumes a COLMAP reconstruction and ground-truth poses per training scene — available for curated data, lacking exactly where robustness is needed. madweight needs neither: it recomputes its outlier test per-scene from reprojection statistics, so it is unaffected by domain shift wherever a partial reconstruction exists.

Measuring the label ceiling. We make the second argument quantitative: on 18 scenes spanning the contamination axis we regenerate labels as one must for a new dataset (from-scratch COLMAP on the same tracks, known intrinsics, no GT poses) and score them against GT-pose reference labels (outlier = positive; coverage = fraction of observations receiving any verdict; ≤ 120 cameras per scene); Table 3. Precision stays high throughout, but recall collapses from 0.82 to 0.26–0.39 and coverage halves once contamination exceeds $\sim 15\%$: the corrupted reconstruction absorbs enough outliers that most reproject acceptably under its distorted geometry, and half the observations receive no label at all. A head trained on such labels is actively taught that most true outliers are inliers, precisely where robustness matters. Conversely, below $\sim 15\%$ COLMAP-derived labels are excellent, consistent with RESfM’s strong in-distribution result and delimiting where the supervised recipe is sound.

5.4. Extreme contamination: the 1DSfM-hard set

To probe beyond 1DSfM’s 43%, we evaluate on the five held-out 1DSfM-hard scenes ($\sim 60\%$ measured contamination, 300 cameras each), adding COLMAP (incremental mapping on the same tracks, known intrinsics) and the full mechanism grid (supplementary table). The ordering extrapolates: over seeds 20–24,

Table 2. Main controlled comparison (RESfM-aligned schedule). Mean / median post-BA translation error. All rows share tracks, architecture, schedule, robust BA, and seeds; only the outlier mechanism and its supervision differ. Upper block = SS-RESfM arms (five-training-seed bands, mean $\pm$ std of per-seed dataset means), as is the supervised baseline; reference and classical rows are single runs, except [§]RESfM (released): a band over five evaluation seeds on the authors’ fixed weights (network held fixed, spread ≤ 0.008 — evaluation noise is negligible; training seeds are what matter). Best per column in bold. ESfM* = vanilla ESfM on ground-truth-cleaned tracks (clean-track reference; mean failure-scene-dominated, as in RESfM’s own Table 1).

Arm	MegaDepth (25.4) $N_c=13126$	1DSfM (43.3) 4563	1DSfM-hard (60) 1497	Strecha (1.7) 54	BMVS (3.1) 225	Olsson (0.5) 3858
Translation error (mean / median), the primary metric						
madweight	0.51 $\pm$ 0.08/0.16	11.0 $\pm$ 1.7/4.9	18.9 $\pm$ 4.1/13.4	2.65 $\pm$ 0.25/1.16	0.14 $\pm$ 0.02/0.05	3.0 $\pm$ 0.3/1.4
weight	0.51 $\pm$ 0.09/0.20	14.4 $\pm$ 3.0/9.6	22.9 $\pm$ 2.8/15.6	1.97 $\pm$ 0.09/0.04	0.14 $\pm$ 0.03/0.02	2.9 $\pm$ 0.3/1.1
weight+TTT	0.40 $\pm$ 0.11/0.18	15.7 $\pm$ 2.1/10.8	22.2 $\pm$ 3.2/15.3	2.00 $\pm$ 0.05/0.12	0.14 $\pm$ 0.02/ 0.01	8.0 $\pm$ 0.3/0.6
remove	0.51 $\pm$ 0.08/0.13	9.1 $\pm$ 2.2/ 1.7 $\pm$ 0.2	12.4 $\pm$ 2.1/10.8	3.20 $\pm$ 0.11/2.03	0.24 $\pm$ 0.06/0.17	3.4 $\pm$ 0.6/1.7
remove+TTT	0.56 $\pm$ 0.07/0.14	8.9 $\pm$ 1.2/2.0	15.2 $\pm$ 1.0/11.3	3.01 $\pm$ 0.07/1.70	0.22 $\pm$ 0.10/0.18	3.6 $\pm$ 0.4/1.9
RESfM-scratch (supervised)	0.37 $\pm$ 0.12/0.07	10.5 $\pm$ 1.5/3.4	19.2 $\pm$ 3.4/16.5	0.39 $\pm$ 0.72/0.01	0.35 $\pm$ 0.04/0.33	7.4 $\pm$ 2.6/ 0.95
RESfM (released) [§]	0.203 $\pm$ 0.008	10.71 $\pm$ 0.27	15.29	0.20 $\pm$ 0.00	0.33 $\pm$ 0.00	9.10
RESfM (paper)	0.169	—	—	—	—	—
RESfM-faithful (authors’ ckpt selection)	0.496	9.75	17.58	0.144	0.367	8.77
ESfM* (clean tracks)	0.60 $\pm$ 0.18	13.24 $\pm$ 0.52	18.12 $\pm$ 4.35	1.14 $\pm$ 0.45	3.79 $\pm$ 3.51	—
ESfM (same tracks, no outlier mech.)	0.71 $\pm$ 0.11	18.42 $\pm$ 0.75	21.76 $\pm$ 0.80	2.03 $\pm$ 0.14	6.45 $\pm$ 3.75	5.91 $\pm$ 3.32
GASfM (released ckpt, our robust BA)	2.25	25.84	36.49	1.14	11.13	4.45
GLOMAP (classical, global)	3.33	27.21	35.42	0.047	0.339	3.17
COLMAP (classical, increm.)	—	—	20.4 $\pm$ 1.1	—	—	—

Buddah Tooth Relic Temple (Olsson, 0.5% contamination), both trained on MegaDepth: the clean-ODD collapse




Figure 3. The clean-ODD collapse, qualitatively (Buddah Tooth Relic Temple, Olsson, 0.5% contamination; both networks trained on MegaDepth, seed 20): top-view camera centers in the evaluation frame; whiskers join each reconstructed camera to its ground-truth position. The label-free madweight arm (left) lands on the ground-truth layout (translation median 0.06, rotation 1.0°); the supervised baseline (right) registers the same number of cameras but collapses them toward a line with large systematic displacements (median 2.55, rotation 16.5°) — on a scene with almost no outliers to reject.

Table 3. The label ceiling: quality of COLMAP-regenerated supervision vs. contamination.

Contam.	n	Cover%	Prec.	Recall	COLMAP t -err
0–5%	6	99.4	0.76	0.82	0.004
5–15%	4	89.6	0.90	0.83	0.014
15–30%	2	71.7	0.83	0.26	1.02
30–45%	3	55.7	0.74	0.39	0.82
45–60%	3	55.1	0.87	0.39	11.0

plain remove attains 12.4 ± 2.1 translation, below the supervised band (19.2 ± 3.4) on every seed pair (medians 10.8 vs. 16.5); madweight (18.9 ± 4.1) statis-

tically matches the baseline; every other mechanism (soft, hybrid, alternative robust statistics) is at 16.9 or worse at seed 20; and COLMAP loses (20.9) even though its error is computed only over the 64% of cameras it registers. Errors are large for every method — these scenes are hard — but the trend matches the label-ceiling prediction: the supervised advantage vanishes exactly where supervision can no longer be manufactured. Strikingly, plain remove on the contaminated tracks also beats ESfM* on oracle-cleaned tracks (12.4 ± 2.1 vs. 18.12 ± 4.35 ; 23 of 25 seed pairs): at 60% contamination, even ground-truth labels used for hard removal destroy more than they fix — deleting

60% of observations starves tracks below the multi-view minimum and thins the viewing graph, whereas adaptive removal with soft-weighted survivors preserves the constraint structure. Having perfect labels would not rescue the removal-by-oracle recipe; the label ceiling is thus not the only obstacle at the extreme — how labels are used matters as much as having them.

5.5. Training on the contaminated domains: the label ceiling is causal

The label-ceiling analysis predicts that supervision fails at high contamination because the labels themselves become untrustworthy — not because of domain shift. We test this causally by making the OOD domains in-distribution: for each dataset we split its scenes into train/validation/held-out test and retrain both recipes on the domain itself (identical protocol; supervised labels generated at reference quality, by the same Appendix-C procedure as MegaDepth’s but against each dataset’s real GT poses — the best case supervision could ever obtain). Table 4 reports five-seed bands on the held-out scenes.

Four findings (Fig. 4). (i) At 60%, supervision loses on its own training domain: the supervised head, trained on these scenes with reference-quality labels, is beaten $\sim 2\times$ at band level by madweight (12.1 ± 6.1 vs. 23.4 ± 7.8 ; four of five madweight seeds beat every supervised seed; medians 9.8 vs. 20.6). “Just train the classifier there” does not work: the high-contamination failure is caused by labeling at contamination, not by domain shift. Sharper still: the supervised head with best-case labels performs identically to no outlier mechanism at all at this contamination (vanilla ESfM trained in-domain: 23.14 ± 5.89 vs. supervised 23.4 ± 7.8) — at 60%, GT-supervised classification adds nothing over doing nothing, while label-free removal halves both. The vanilla row completes a symmetric statement: GT-label supervision matches no mechanism at all at both extremes of the axis — identical at 60% (23.4 vs. 23.1) and overlapping at 0.5% in-domain (2.03 ± 1.86 vs. 1.45 ± 1.44) — earning its keep only in the curated middle band (25–43%). The absolute errors deserve decomposition (supplementary figure): they are not uniformly “large for everyone.” On Roman_Forum, madweight produces usable reconstructions (median 1.6–3.5, rotation 2.3–8.2°, in four of five seeds) where supervision fails outright (medians 9.5–36); on Union_Square the supervised errors look lower but are computed over only 43–55 of 300 registered cameras, versus madweight’s 137–186 — the error-vs-coverage trade-off of §4.2 again, with the label-free arm covering 3–4 $\times$ more of the scene at moderate error. (ii) The BlendedMVS failure was never domain

shift either: supervision trained on BlendedMVS lands at 0.20–0.37 in four of five seeds — indistinguishable from the MegaDepth-trained baseline evaluated OOD (0.35 ± 0.04) — with a single escape seed (0.02), the same bimodality as its clean-OOD behaviour, while label-free weighting is seed-stable at 0.114 ± 0.006 . Class imbalance alone cannot explain this: at 0.5% (Olsson), where positives are rarer still, supervision trained in-domain reaches 0.41–1.08 in three of five seeds — extreme positive scarcity does not preclude learning — whereas BlendedMVS never reaches that level on any seed; the BlendedMVS failure is a dataset-specific dead zone (its rotation pathology is the leading suspect), not a general low-contamination bound. (iii) In-domain training attenuates but does not cure the clean-benchmark unreliability: the recipe that lands at 7.4–9.1 on Olsson MegaDepth-trained improves to 2.03 ± 1.86 trained in-domain — but bimodally, excellent in three seeds (0.41–1.08) and degraded in two (3.9–4.2, rotation spread $12 \pm 12^\circ$). The OOD collapse is thus mostly the MegaDepth-fit head misclassifying on clean scenes (in-domain labels remove the bulk of it), yet the seed-reliability asymmetry survives even here: weight+TTT (1.60 ± 0.81) leads the column, madweight (1.97 ± 0.94) matches the supervised mean at half the spread, and no label-free arm has a failure seed. Supervision’s clean-data advantage is a three-in-five-seeds advantage, not a dependable one. (iv) The supervised advantage inverts between 43% and 60% (4.97 ± 2.08 vs. madweight 6.37 ± 0.20 at 43%; a $2\times$ loss at 60%), and the complementarity ordering replicates in-distribution — the mechanism flip is a property of the contamination level, not the train/test relationship. Read as a curve, best-case-label supervision wins at 25% and narrowly 43%, is seed-bimodal at 0.5% (tied in mean, far wider in spread), and loses at 3.1% and 60%: two distinct failure modes — a dataset-specific dead zone, and the causal label ceiling that no in-domain retraining can cross — with seed-unreliability shadowing the whole axis.

6. Limitations

(1) No single deployable configuration wins all datasets; selecting the best mechanism per dataset (Table 2) requires knowing the contamination level. We present a study of the phenomenon, not a turnkey adaptive method; an unsupervised selector is future work. (2) Supervised RESfM retains in-distribution MegaDepth (the self-supervised arms come within $\sim 38\%$ at the band level, no closer) and the clean Strecha mean, and RESfM-released remains stronger on MegaDepth, so we claim no SOTA. (3) madweight matches rather than beats supervision at mod-

Table 4. In-distribution comparison on the (formerly OOD) domains, ordered by contamination: mean translation over held-out test scenes, mean $\pm$ std over five training seeds per recipe. Supervised labels are GT-derived (best case). Best per column in bold.

Arm	Olsson @0.5%	BlendedMVS @3.1%	1DSfM @43%	1DSfM-hard @60%
madweight	1.97 $\pm$ 0.94	0.093 $\pm$ 0.099	6.37 $\pm$ 0.20	12.1 $\pm$ 6.1
weight	3.05 $\pm$ 1.39	0.114 $\pm$ 0.006	10.65 $\pm$ 2.49	28.1 $\pm$ 5.9
weight+TTT	1.60 $\pm$ 0.81	0.081 $\pm$ 0.046	9.91 $\pm$ 2.39	29.3 $\pm$ 8.4
remove	2.94 $\pm$ 0.27	0.228 $\pm$ 0.080	6.68 $\pm$ 1.44	19.4 $\pm$ 7.7
RESfM (sup., GT labels)	2.03 $\pm$ 1.86	0.260 $\pm$ 0.149	4.97 $\pm$ 2.08	23.4 $\pm$ 7.8
ESfM (no mechanism)	1.45 $\pm$ 1.44	2.36 $\pm$ 3.27	10.20 $\pm$ 2.97	23.14 $\pm$ 5.89




Figure 4. The completed in-distribution contamination curve. Post-BA mean translation (log scale) when training and test share the domain, versus measured track contamination. The supervised recipe uses best-case labels (GT-derived; on MegaDepth, its native COLMAP labels). Each label-free arm is its own line: soft weighting ($\pm$ TTT) leads at low contamination, robust removal (madweight, remove) at high — the complementarity mechanism flip replicated in-distribution. Supervision wins at 25% and narrowly at 43%, is seed-bimodal at 0.5% (mean tied with the label-free arms at more than twice their spread), and loses at 3.1% (dataset-specific dead zone) and 60% (the causal label ceiling): its failures are two distinct modes, not a smooth function of contamination. The rotation panel (right) tells the same story and makes the BlendedMVS dead zone starker still (supervised $20.7 \pm 12.2^\circ$ vs. label-free $4\text{--}13^\circ$). Bands: $\pm$ std over five training seeds (MegaDepth rotation: five-seed means, bands omitted).

erate/extreme contamination; plain remove leads every baseline seed pairwise at $\sim 60\%$ and on medians at 43%, but over a reduced registered set (§4.2), and the magnitude of the supervised deficit at BlendedMVS and high contamination is learning-rate-dependent (supplementary) — the Olsson collapse and the mechanism flip are not. (4) OOD tracks are self-constructed (not comparable to RESfM’s published OOD numbers). (5) Two symmetric-structure scenes fail for all methods — the cases where excess removal disconnects the viewing graph [1].

7. Conclusion

SS-RESfM makes deep SfM outlier handling entirely self-supervised while staying within $\sim 38\%$ of the supervised baseline in distribution, winning two of three clean-OOD splits — reliably, on every training seed — and beating it at high contamination on every seed pair, where both beat COLMAP and supervision itself can no longer be manufactured. The optimal

mechanism is contamination-dependent: removal at high contamination, reweighting at low. Robust deep SfM should treat outlier handling as adaptive and self-supervised; code, configurations, and the rebuilt OOD tracks will be released upon publication.

References

- [1] Khatib, Kasten, Moran, Galun, and Basri. RESfM: Robust Deep Equivariant Structure from Motion. ICLR 2025. 1, 2, 3, 9
- [2] Moran, Koslowsky, Kasten, Maron, Galun, and Basri. Deep Permutation Equivariant Structure from Motion. ICCV 2021. 1, 2, 3
- [3] Brynte, Iglesias, Olsson, and Kahl. Learning Structure-from-Motion with Graph Attention Networks. CVPR 2024. 2
- [4] Schönberger and Frahm. Structure-from-Motion Revisited. CVPR 2016. 1
- [5] Rousseeuw and Croux. Alternatives to the Median Absolute Deviation. JASA 1993. 3

- 653 [6] Wilson and Snavely. Robust Global Translations with
654 1DSfM. ECCV 2014. 3
- 655 [7] Li and Snavely. MegaDepth: Learning Single-View
656 Depth Prediction from Internet Photos. CVPR 2018.
657 3
- 658 [8] Strecha, von Hansen, Van Gool, Fua, and Thoen-
659 nessen. On Benchmarking Camera Calibration and
660 Multi-View Stereo for High-Resolution Imagery.
661 CVPR 2008. 3
- 662 [9] Yao et al. BlendedMVS: A Large-Scale Dataset
663 for Generalized Multi-View Stereo Networks. CVPR
664 2020. 3
- 665 [10] Khatib, Galun, and Basri. Learning Global Camera
666 Poses from Noisy View-Graphs for Structure from Mo-
667 tion (VGPA). ECCV 2026. 2
- 668 [11] Olsson and Enqvist. Stable Structure from Motion for
669 Unordered Image Collections. SCIA 2011. 3