

Supplementary Material: Environment Sensitivity and Reproducibility of Post-BA SfM Benchmarks

Anonymous 3DV submission

Paper ID 524

A. What this document supports

This supplementary substantiates four specific statements in the main paper; a reviewer following any of its pointers can go directly to the listed section.

- Protocol, “Track provenance”: “ ~ 0.209 is the reproducible level [for the released RESfM checkpoint] ... unchanged (0.217) under the authors’ exact pinned software stack” (Secs. B–D).
- Protocol: “a stack-replication study (supplementary) shows the no-TTT arms are insensitive to it, while test-time fine-tuned means inherit solver-level nondeterminism” (Sec. E).
- Table 2, † footnote: “TTT is run-to-run nondeterministic” (Sec. E).
- Reading any post-BA mean in this literature: differences $\lesssim 0.05$ are within environment noise (Sec. F).

In addition, Sec. G tabulates the per-scene translation, rotation, and registered-camera (N_r) results behind the main paper’s dataset-level numbers, with dataset-level rotation and registration summaries in Table 11; Sec. H gives the complete training and evaluation configuration, Sec. I provides the per-scene label-degradation breakdown behind the label-ceiling analysis, and Sec. J provides extended versions of the main paper’s Tables 2, 5, 6, and 7 (adding rotation and N_r) intended to replace them in the camera-ready.

B. The reproduction question

Evaluating the RESfM authors’ released checkpoint in our harness on the 36 MegaDepth test scenes yields a post-BA mean translation error of 0.209 vs. the published 0.169 (rotation 1.89° vs. 1.29°). Since our controlled comparisons hinge on baseline fidelity, we traced this discrepancy to its source. The finding is of independent interest: post-BA means on internet-photo benchmarks carry a practical environment-noise floor of $\mathcal{O}(0.05)$, produced by chaotic basin selection in robust bundle adjustment.

Table 8. Released RESfM checkpoint, MegaDepth 36-scene post-BA means, under three conditions, vs. the published number. N_r is the mean number of registered cameras per scene; the published row averages the paper’s per-scene Table 1 values.

Condition	Trans	Rot (°)	N_r
Published (paper Table 1)	0.169	1.29	298
Our env (pip/OpenBLAS), 5-seed mean	0.203	1.82	304
Our env, seed 20	0.209	1.89	304
Authors’ exact stack (conda/MKL)	0.217	1.92	304

C. Eliminating the controllable factors

Data. Our MegaDepth npz files are byte-identical to the authors’ released archive, including the Group-2 300-camera subsamples.

Randomness. Five evaluation seeds vary the pre-BA network output by ± 0.004 and the post-BA mean by ± 0.01 ; negligible.

Network forward pass. The pre-BA output (translation/rotation $\sim 1.73/11.9^\circ$) is essentially seed- and environment-stable; the divergence therefore enters at the robust-BA stage. (Notably, the paper does not report pre-BA numbers; all its tables are post-robust-BA.)

Software stack. We rebuilt the authors’ exact pinned environment from their released environment.yaml: Python 3.8.20, MKL-backed numpy 1.24.3 (mkl=2023.1.0), pycolmap 3.10.0, pyceres 2.3, lightning 2.3.3. (The sole deviation is the CUDA build of PyTorch, forced by GPU-driver compatibility; the BA stage is CPU-only.) Re-evaluating all 36 scenes in this matched environment gives translation/rotation errors of 0.217/1.92°, statistically indistinguishable from our environment’s 0.209/1.89°, and not the published 0.169/1.29°. Registration is equally stable: all three of our conditions register 304 cameras per scene on average (Table 8), slightly more than the published 298, so the higher mean error is not an artifact of evaluating fewer cameras.

066	D. Mechanism: chaotic basin selection in ro-	F. Implications for reading the paper's num-	116
067	bust BA	bers	117
068	Initialization perturbation is at the 10^{-10} level. We	(1) The published number and ours are draws from	118
069	isolated the BLAS layer by running the identical DLT	a distribution over environments; no software match-	119
070	triangulation (the BA initialization) on identical inputs	ing can force a particular draw. (2) Differences	120
071	under OpenBLAS and MKL, single-threaded: over	of $\lesssim 0.05$ in post-BA translation means on these	121
072	11,131 triangulated points the maximum relative dif-	benchmarks are within environment noise and should	122
073	ference is 4.8×10^{-10} (mean 5×10^{-14} ; NaN patterns	not be read as method differences, by any paper.	123
074	identical).	(3) Consequently, meaningful comparisons must be	124
075	Yet outcomes move by $\mathcal{O}(0.1)$. This sub-nanoscale per-	environment-controlled: every number in the main pa-	125
076	turbation changes the per-scene post-BA translation by	per (all arms, RESfM-from-scratch, and the released	126
077	up to 0.167, with 16/36 scenes moving by more than	checkpoint) is produced in a single environment, and	127
078	0.01. Robust BA on ill-conditioned scenes amplifies	the out-of-distribution margins ($2\times$ and larger) sit far	128
079	last-bit differences into different convergence basins.	above this noise floor.	129
080	The published gap has the same signature. Compar-	G. Per-scene results	130
081	ing our released-checkpoint evaluation to the paper's	Tables 9 and 10 report the per-scene post-BA transla-	131
082	Table 1 per scene: 20/36 scenes agree within 0.05 (8	tion and rotation errors and registered-camera counts	132
083	essentially exactly), 17 are better in our runs, and the	underlying the main paper's dataset-level numbers	133
084	entire mean gap is carried by four basin flips (scenes	(seed 20, the declared pipeline): the 36 MegaDepth test	134
085	0046, 0148, 0012, 5016; e.g. 0046: $0.043 \rightarrow 0.892$).	scenes and the ten 1DSfM scenes, for the three self-	135
086	Bidirectional, scene-concentrated flips are the signa-	supervised arms (madweight, weight, weight+TTT)	136
087	ture of basin selection, not of a weaker checkpoint,	and the supervised RESfM baseline trained from	137
088	which would shift errors in one direction rather than	scratch under the identical recipe. The per-scene	138
089	improve 17 of 36 scenes.	view makes two aggregate statements concrete: on	139
090	E. Stack replication and TTT run-to-run vari-	1DSfM no single arm dominates, the best mech-	140
091	ance	anism flips scene by scene (with <u>Ellis_Island</u> and	141
092	Beyond the released checkpoint, we re-evaluated both	<u>Tower_of_London</u> hard for every arm), consistent	142
093	from-scratch main-comparison checkpoints under the	with the per-scene complementarity analysis in the	143
094	authors' pinned stack (MKL numpy, pycolmap 3.10.0,	main paper; on MegaDepth the four arms are close	144
095	pyceres 2.3), changing nothing else, and additionally	on most scenes and the dataset-level gap is dominated	145
096	repeated the TTT evaluation within our own stack.	by a handful of high-error scenes. Per scene, the	146
097	The supervised baseline is essentially unchanged across	catastrophic scenes (<u>Ellis_Island</u> , <u>Tower_of_London</u>)	147
098	stacks: MegaDepth $0.344 \rightarrow 0.350$ mean, $0.070 \rightarrow$	fail in both metrics, though rotation does not always	148
099	0.088 median, only 7/36 scenes moving more than 0.05	track translation (weight's <u>Vienna_Cathedral</u> failure	149
100	(max ± 0.37 , bidirectional). This matches the released-	is translation-specific: $22.71/0.46$), and N_r stays in	150
101	checkpoint result above. The SS-RESfM weight+TTT	the 58–92% range of each scene's images on 1DSfM	151
102	arm, by contrast, varies regardless of stack: three eval-	with no arm systematically dropping scenes. Table 11	152
103	uations of the same checkpoint (ours, a same-stack re-	completes the dataset-level record with rotation and	153
104	peat, the authors' stack) give MegaDepth means 0.316,	registered-camera counts alongside translation. Ro-	154
105	0.401, 0.453 (~ 18 – $20/36$ scenes shifting per repeat, up	tation broadly agrees with translation at the dataset	155
106	to ± 2.9 , largely the same scenes, medians 0.12–0.19).	level: the supervised BlendedMVS failure is even larger	156
107	The attribution is therefore not the software stack but	in rotation (36°) and the high-contamination order-	157
108	test-time fine-tuning itself: it runs an optimization at	ing is preserved, though the best arm can differ per	158
109	evaluation time, so the chaotic basin selection docu-	metric (on MegaDepth, weight+TTT leads transla-	159
110	mented above re-enters the result on every run. Con-	tion while RESfM leads rotation). Registration is	160
111	sequently the paper reports all results from the sin-	contamination-dependent: the supervised baseline reg-	161
112	gle declared pipeline, treats TTT means as carrying	isters the most cameras in distribution (85% vs. 78–	162
113	$\sim \pm 0.07$ run-to-run variance, and rests its environ-	80% on MegaDepth) and on clean Strecha, while un-	163
114	ment- and run-independent claims on the deterministic no-	der high contamination the label-free arms retain more	164
115	TTT cells (which reproduce exactly on repeat).	cameras than supervision (38–43% vs. 37% on 1DSfM-	165
		hard) at lower error (madweight).	166

Table 9. Per-scene MegaDepth results (layout follows RESfM’s Table 1). For each scene we show the number of input images (N_c) and the fraction of outlier observations; for each model, the number of registered cameras (N_r) and mean rotation (e_R , degrees) and translation (e_t) errors, post-BA, seed 20. Above the middle rule are Group 1 scenes (< 1000 images); below are Group 2 scenes (> 1000 images, subsampled to 300 for testing), each group ordered by outlier fraction. Winning results per scene are marked in bold and underlined. weight+TTT values are single draws of a run-to-run nondeterministic evaluation (Sec. E).

Scene	N_c	%out	madweight			weight			weight+TTT			RESfM		
			N_r	e_R	e_t	N_r	e_R	e_t	N_r	e_R	e_t	N_r	e_R	e_t
0238	522	45%	437	19.02	1.41	137	8.23	0.62	141	8.37	0.63	355	<u>1.40</u>	<u>0.12</u>
0060	528	42%	469	0.29	0.03	476	0.20	0.02	452	4.13	0.37	494	<u>0.04</u>	<u>0.01</u>
0197	870	41%	470	7.22	0.89	341	3.54	0.38	422	6.94	1.11	739	<u>0.97</u>	<u>0.11</u>
0094	763	40%	493	2.72	0.45	505	1.44	<u>0.16</u>	432	<u>1.34</u>	0.17	572	5.33	0.88
0265	571	39%	457	17.62	3.26	477	5.78	2.34	547	<u>0.60</u>	<u>0.19</u>	424	16.27	3.09
0083	635	31%	342	4.34	0.23	511	17.70	1.56	317	3.20	0.20	582	<u>1.40</u>	<u>0.10</u>
0076	558	30%	423	1.40	0.38	519	0.61	0.23	454	0.74	0.22	499	<u>0.25</u>	<u>0.02</u>
0185	368	30%	347	0.21	0.01	335	0.24	0.04	321	0.88	0.09	352	<u>0.07</u>	<u>0.01</u>
0048	512	24%	475	1.93	0.14	467	<u>0.35</u>	<u>0.02</u>	450	0.43	0.03	489	1.85	0.11
0024	356	23%	299	1.27	0.18	291	<u>1.13</u>	<u>0.17</u>	289	4.78	1.13	334	3.67	1.09
0223	214	17%	205	2.66	0.24	198	1.69	0.23	199	3.06	0.30	204	<u>0.45</u>	<u>0.08</u>
5016	28	17%	23	0.27	0.04	25	0.08	0.02	28	<u>0.07</u>	0.02	28	0.08	<u>0.01</u>
0046	440	15%	422	<u>2.88</u>	<u>0.22</u>	421	4.11	0.30	419	3.88	0.25	410	14.64	0.93
0099	299	47%	224	<u>0.81</u>	0.15	216	2.77	3.52	212	1.58	<u>0.13</u>	252	2.18	0.39
1001	285	44%	280	1.58	1.22	278	1.49	0.73	279	<u>1.22</u>	<u>0.49</u>	273	1.23	0.50
0231	296	42%	274	1.63	0.15	257	0.64	0.07	264	1.49	0.24	265	<u>0.17</u>	<u>0.02</u>
0411	299	30%	281	0.22	0.03	245	2.99	0.32	250	4.12	0.46	275	<u>0.11</u>	<u>0.02</u>
0377	295	28%	200	0.62	0.05	222	2.76	0.26	243	<u>0.56</u>	<u>0.03</u>	223	0.82	0.06
0102	299	26%	289	0.15	0.02	289	<u>0.07</u>	<u>0.01</u>	291	0.08	0.01	287	0.10	0.02
0147	298	25%	261	31.69	2.99	235	15.28	1.21	225	15.27	<u>1.20</u>	247	<u>4.61</u>	3.22
0148	287	25%	220	2.60	0.33	201	2.11	0.11	196	<u>0.62</u>	0.04	202	0.94	<u>0.04</u>
0446	298	22%	292	<u>0.40</u>	<u>0.03</u>	285	1.20	0.19	292	1.04	0.19	293	0.71	0.05
0022	297	21%	275	0.14	0.02	272	0.22	0.03	267	0.33	0.04	276	<u>0.11</u>	<u>0.02</u>
0327	298	21%	282	1.97	0.09	285	0.22	0.02	284	0.21	0.02	277	<u>0.09</u>	<u>0.01</u>
0015	284	21%	196	2.64	0.30	196	3.87	0.56	211	3.90	0.65	232	<u>1.31</u>	<u>0.30</u>
0455	298	20%	283	<u>0.20</u>	<u>0.01</u>	285	0.21	0.01	288	0.46	1.60	293	0.85	0.09
0496	297	19%	279	<u>0.27</u>	0.04	285	0.50	<u>0.04</u>	287	1.28	0.09	283	0.46	0.07
1589	299	17%	265	0.96	0.06	262	0.28	0.05	296	1.91	0.30	291	<u>0.06</u>	<u>0.01</u>
0012	299	16%	281	<u>0.23</u>	<u>0.02</u>	247	9.51	0.70	262	2.76	0.22	288	5.53	0.34
0104	284	16%	189	0.32	0.03	193	2.28	0.18	194	1.44	0.07	199	<u>0.30</u>	<u>0.02</u>
0019	299	15%	232	3.22	0.17	253	6.55	0.28	264	3.24	0.13	252	<u>0.05</u>	<u>0.01</u>
0063	293	14%	258	0.63	0.08	268	0.72	0.07	267	0.65	<u>0.06</u>	266	<u>0.44</u>	0.07
0130	285	14%	250	10.17	1.60	184	1.87	0.16	185	<u>0.33</u>	<u>0.01</u>	195	1.03	0.07
0080	284	13%	132	1.10	<u>0.05</u>	148	2.48	0.20	147	0.70	0.07	141	<u>0.66</u>	0.10
0240	298	12%	275	<u>2.23</u>	<u>0.25</u>	275	4.64	0.34	272	5.32	0.46	282	3.02	0.35
0007	290	12%	106	<u>0.07</u>	<u>0.01</u>	269	1.86	0.09	261	3.28	0.14	104	0.99	0.04

H. Implementation details

Table 12 lists the full training and evaluation configuration. Both methods share every setting except the base learning rate (10^{-4} self-supervised, 10^{-3} supervised, each selected on its own training objective) and the supervision source (COLMAP-derived labels vs. the adaptive percentile pseudo-labels).

Table 10. Per-scene 1DSfM results (layout follows RESfM’s Table 1). Columns as in Table 9; seed 20, scenes ordered by outlier fraction. Winning results per scene are marked in bold and underlined.

Scene	N_c	%out	madweight			weight			weight+TTT			RESfM		
			N_r	e_R	e_t	N_r	e_R	e_t	N_r	e_R	e_t	N_r	e_R	e_t
Ellis_Island	227	56%	180	<u>17.16</u>	17.59	171	17.37	<u>15.90</u>	177	17.74	17.11	192	18.15	16.02
NYC_Library	332	54%	282	2.32	2.94	261	<u>1.16</u>	<u>1.19</u>	279	10.30	3.39	285	12.99	6.28
Vienna_Cathedral	836	51%	537	7.42	7.30	484	<u>0.46</u>	22.71	651	26.80	33.57	492	1.09	<u>2.29</u>
Notre_Dame	553	48%	507	3.33	0.68	491	1.00	0.34	485	<u>0.83</u>	<u>0.33</u>	493	1.08	0.49
Yorkminster	437	45%	299	8.05	<u>5.79</u>	371	14.95	30.87	380	12.20	10.79	309	<u>6.88</u>	8.77
Piazza_del_Popolo	338	42%	218	<u>0.43</u>	<u>0.94</u>	265	19.54	10.66	254	20.46	17.05	255	3.15	2.49
Tower_of_London	472	42%	317	35.61	76.34	337	23.30	<u>28.29</u>	314	<u>17.91</u>	72.61	299	49.66	67.95
Montreal_Notre_Dame	450	36%	311	0.59	0.94	325	<u>0.55</u>	11.35	320	0.81	1.24	400	0.59	<u>0.93</u>
Alamo	577	29%	467	2.45	1.77	507	2.74	2.18	497	2.49	<u>1.67</u>	502	<u>2.42</u>	1.73
Madrid_Metropolis	341	29%	249	5.72	9.63	301	2.84	4.37	295	2.35	3.52	287	<u>1.82</u>	<u>3.38</u>

Table 11. Dataset-level translation, rotation, and registered cameras (post-BA, seed 20) for the four main-comparison arms. Per-scene values are camera means; N_r is the fraction of ground-truth cameras surviving robust BA and the 5-pixel reprojection quality mask, summed over the dataset’s scenes.

	madweight	weight	weight+TTT	RESfM
Translation error (mean / median over scenes)				
MegaDepth	0.422/0.145	0.423/0.182	0.316/0.186	0.344/0.07
1DSfM	12.4/4.37	12.8/11.0	16.1/7.16	11.0/2.93
1DSfM-hard	12.2/8.83	21.9/13.1	26.4/17.7	16.2/13.1
Strecha	2.5/0.97	2/0.026	1.98/0.005	0.006/0.005
BlendedMVS	0.146/0.004	0.1/0.036	0.136/0.011	0.379/0.348
Olsson	2.57/1.15	2.65/1.22	7.74/0.18	8.47/0.17
Rotation error ($^\circ$, mean / median over scenes)				
MegaDepth	3.49/1.33	3.04/1.77	2.51/1.39	2.01/0.84
1DSfM	8.31/4.53	8.39/2.79	11.19/11.25	9.79/2.79
1DSfM-hard	14.40/7.39	25.58/17.70	31.97/12.83	22.70/15.86
Strecha	14.62/8.43	7.48/0.17	7.41/0.02	0.02/0.01
BlendedMVS	9.60/0.38	6.65/3.06	7.91/0.55	36.10/29.55
Olsson	13.57/4.37	13.23/5.41	10.49/2.27	12.34/3.35
Registered cameras N_r				
MegaDepth	80%	79%	78%	85%
1DSfM	74%	77%	80%	77%
1DSfM-hard	38%	39%	43%	37%
Strecha	91%	96%	98%	100%
BlendedMVS	95%	88%	95%	90%
Olsson	83%	90%	90%	88%

Table 12. Training and evaluation configuration. All arms in the main paper use these settings; the only per-method difference is the base learning rate.

Setting	Value
Training (both methods)	
Epochs	20,000 (one epoch = one pass over all scenes)
Batch	4 scenes
LR schedule	milestones 10k, 11k, $\dots$, 20k, $\gamma = 0.9$
Base LR	10^{-4} (SS-RESfM) / 10^{-3} (RESfM)
Loss weights	$\alpha_{\text{reproj}} = 1$, $\beta_{\text{cls}} = 1$
Infinity-point margin	10^{-4}
Checkpoint selection	best validation error, evaluated every 500 epochs
Seed	20 (21–22 / 21–24 for the seed studies)
Adaptive pseudo-labels (SS-RESfM only)	
Inlier / outlier percentile	20 / 80 (of per-observation reproj. error)
Threshold separation	scale-aware floor, $0.25 \times$ median error
Warm-up	10 epochs before pseudo-labeling starts
Evaluation	
Removal threshold τ	0.6 (0.8 for low-contamination single-camera sets, following RESfM)
Test-time fine-tuning	1,001 iterations, LR 5×10^{-3} , evaluated every 250 iterations
Bundle adjustment	pyceres/pycolmap robust BA, as RESfM

Table 13. Per-scene label-degradation results. For each scene: contamination (fraction of track observations that are GT outliers), COLMAP-label coverage (fraction of observations receiving a label), and precision/recall of the COLMAP-derived outlier labels against GT, sorted by contamination. Above $\sim 15\%$ contamination, label recall never exceeds 0.50: the labels miss at least half of the true outliers.

Dataset	Scene	Contam. (%)	Coverage	Prec.	Recall
BlendedMVS	58c4bb4f4a...	0.5	1.00	0.86	0.94
Strecha	Herz-Jesu-P8	0.7	1.00	0.51	0.66
Strecha	fountain-P11	0.7	1.00	0.69	0.76
Strecha	Herz-Jesu-P25	2.6	0.99	0.72	0.82
Strecha	entry-P10	2.7	0.99	0.88	0.88
BlendedMVS	5be883a4f9...	4.0	0.98	0.90	0.84
BlendedMVS	5b6e716d67...	6.1	0.92	0.89	0.92
BlendedMVS	5bfd0f32ec...	7.1	0.97	0.97	0.97
BlendedMVS	5b950c7160...	9.2	0.96	0.92	0.94
MegaDepth	0007	11.2	0.73	0.84	0.47
1DSfM	Alamo	26.6	0.76	0.81	0.25
1DSfM	Madrid_Metropolis	28.8	0.67	0.85	0.28
MegaDepth	0041	31.4	0.49	0.80	0.42
MegaDepth	0060	39.0	0.42	0.53	0.45
1DSfM	Piazza_del_Popolo	41.9	0.76	0.90	0.30
MegaDepth	0047	54.9	0.37	0.83	0.50
1DSfM	Ellis_Island	57.9	0.79	0.89	0.37
1DSfM-hard	Gendarmenmarkt	59.8	0.50	0.89	0.29

I. Per-scene label degradation

Table 13 gives the per-scene data behind the main paper’s label-ceiling analysis: for each of the 18 scenes we measure the contamination level, the fraction of track observations that receive a COLMAP-derived label at all (coverage), and the precision and recall of those labels against ground truth. The aggregate trend in the main paper is visible scene by scene: below $\sim 10\%$ contamination coverage is 0.92–1.00 and recall is 0.66–0.97 (mostly above 0.8), while above $\sim 15\%$ contamination recall never exceeds 0.50 on any scene (0.25–0.50) and coverage falls as low as 0.37: the supervision signal degrades exactly where outlier handling matters most.

Table 14. Extended main comparison (extends main-paper Table 2; RESfM-aligned schedule, seed 20). Translation panel repeats the main table (mean / median; $^\dagger/^\ddagger$ as there); rotation (mean / median, $^\circ$) and N_r panels at seed 20. Best per column in bold.

Arm	MegaDepth (25.4)	1DSfM (43.3)	1DSfM-hard (60)	Strecha (1.7)	BMVS (3.1)	Olsson (0.5)
Translation error (mean / median), the primary metric						
madweight	0.422/0.145	12.39/4.37	17.0 $\pm$ 4.4/ 11.7‡	2.50/0.97	0.146/ 0.004	2.8 $\pm$ 0.3 /1.2 ‡
weight	0.423/0.182	12.79/11.01	21.9 $\pm$ 3.4/12.7 ‡	2.00/0.026	0.100 /0.036	2.8 $\pm$ 0.4/1.1 ‡
weight+TTT	0.39 $\pm$ 0.07/0.14 †	16.13/7.16	21.9 $\pm$ 3.9/14.5 ‡	1.98/ 0.005	0.136/0.011	8.0 $\pm$ 0.4/0.4 ‡
RESfM-scratch (supervised)	0.344/0.070	11.03/2.93	16.22 /13.08	0.006/0.005	0.379/0.348	8.47/ 0.17
RESfM (released)	0.209	10.71	—	0.197	0.330	—
RESfM (paper)	0.169	—	—	—	—	—
Rotation error ($^\circ$, mean / median)						
madweight	3.49/1.33	8.31 /4.53	14.40/7.39	14.62/8.43	9.60/0.38	13.57/4.37
weight	3.04/1.77	8.39/ 2.79	25.58/17.70	7.48/0.17	6.65 /3.06	13.23/5.41
weight+TTT	2.51/1.39	11.19/11.25	31.97/12.83	7.41/0.02	7.91/ 0.55	10.49/2.27
RESfM-scratch (supervised)	2.01/0.84	9.79/ 2.79	22.70/15.86	0.02/0.01	36.10/29.55	12.34/3.35
Registered cameras N_r						
madweight	80%	74%	38%	91%	95%	83%
weight	79%	77%	39%	96%	88%	90%
weight+TTT	78%	80%	43%	98%	95%	90%
RESfM-scratch (supervised)	85%	77%	37%	100%	90%	88%

J. Extended main-paper tables (camera-ready)

Tables 14–17 extend the main paper’s Tables 2, 5, 6, and 7 with rotation error and registered-camera counts alongside translation, following RESfM’s metric emphasis (translation first, then rotation, then N_r): Table 14 extends the main comparison (Table 2), Table 15 the full mechanism matrix (Table 5), Table 16 the removal-statistic ablation (Table 6), and Table 17 the per-scene selection (Table 7). Table 14 reports the seed-20 evaluations behind the main paper; Tables 15 and 16 report mean $\pm$ std over five independent training seeds (20–24) for every dataset and metric. In the camera-ready we will replace main-paper Tables 2, 5, 6, and 7 with these extended versions. Reading the additions (Table 14): rotation broadly agrees with translation (the supervised BlendedMVS failure is even larger in rotation, 36 $^\circ$; the high-contamination ordering is preserved), though the best arm can differ per metric (on MegaDepth, weight+TTT leads translation while the supervised baseline leads rotation). Registration is contamination-dependent: the supervised baseline registers the most cameras in distribution (85%) and on clean Strecha (100%), while under high contamination the label-free arms retain more (38–43% vs. 37% on 1DSfM-hard). In the ablation (Table 16) rotation and registration are statistic-insensitive on 1DSfM (band means 7.6–9.0 $^\circ$, 70–75%): the statistic matters through translation. The selection oracle (Table 17) improves both metrics on MegaDepth — translation 0.116 vs. 0.422 for the best fixed arm and 0.344 for the supervised RESfM-from-scratch baseline; rotation 1.22 $^\circ$ vs. the baseline’s 2.01 $^\circ$ — but registers fewer 1DSfM cameras (66% vs. 74–77% for the fixed arms): selecting by translation error can sacrifice registration.

Two notions of registration, and an error-vs-coverage trade-off. Two registration notions must be distinguished. Main-paper Table 4’s “Registered” column reports pose coverage: every learned arm outputs a pose for every camera (“all”), whereas COLMAP genuinely fails to register 36% of the cameras. N_r here is stricter: the cameras surviving robust BA and the 5-pixel reprojection quality mask, i.e. the set over which the reported errors are computed. Under this stricter notion (Table 15, N_r panel) the plain-removal arms attain their low errors over a reduced registered set: on 1DSfM plain remove registers 58% vs. the supervised baseline’s 77%, and on 1DSfM-hard 29% vs. 34% (remove+TTT: 58%/32%), consistent with aggressive pruning disconnecting weakly-supported cameras (main-paper limitation 5). The camera-ready will state this caveat alongside the plain-remove claims (“over the cameras it retains”). madweight, by contrast, registers comparably to the supervised baseline on 1DSfM (70% vs. 77%) and more under extreme contamination (43% vs. 34% on 1DSfM-hard), so its statistical match at extreme contamination and its clean-OOD wins are not coverage artifacts, and the soft arms register at least as many cameras as supervision under high contamination. These readings, together with the paragraph above, are intended for the camera-ready main text (the Metrics paragraph, Secs. 5.1 and 5.5, and limitation 5).

Table 15. Extended full shallow mechanism matrix, five-seed bands (extends main-paper Table 5). All entries are mean $\pm$ std over five independent training seeds (20–24); ⁽ⁿ⁾ marks cells currently backed by fewer seeds (remaining seed evaluations in flight). N_r is the percentage of the dataset’s ground-truth cameras surviving robust BA and the 5-pixel quality mask.

Arm	MD	1DSfM	1DSfM-hard	Strecha	BMVS	Olsson
Translation error (mean $\pm$ std over training seeds)						
remove	0.51 $\pm$ 0.08	9.13 $\pm$ 2.17	12.42 $\pm$ 2.05	3.20 $\pm$ 0.11	0.24 $\pm$ 0.06	3.41 $\pm$ 0.57
remove+TTT	0.56 $\pm$ 0.07	8.93 $\pm$ 1.20	15.15 $\pm$ 0.98	3.01 $\pm$ 0.07	0.22 $\pm$ 0.10	3.63 $\pm$ 0.35
hybrid	0.48 $\pm$ 0.10 ⁽³⁾	15.99 $\pm$ 1.36 ⁽³⁾	26.35 ⁽¹⁾	3.09 $\pm$ 0.45 ⁽³⁾	0.13 $\pm$ 0.03 ⁽³⁾	2.82 ⁽¹⁾
hybrid+TTT	0.47 $\pm$ 0.13 ⁽³⁾	15.38 $\pm$ 1.07 ⁽³⁾	24.99 ⁽¹⁾	3.19 $\pm$ 0.31 ⁽³⁾	0.14 $\pm$ 0.02 ⁽³⁾	3.19 ⁽¹⁾
madweight	0.51 $\pm$ 0.08	11.04 $\pm$ 1.68	18.92 $\pm$ 4.05	2.65 $\pm$ 0.25	0.14 $\pm$ 0.02	3.01 $\pm$ 0.34
weight	0.51 $\pm$ 0.09	14.35 $\pm$ 2.97	22.92 $\pm$ 2.79	1.97 $\pm$ 0.09	0.14 $\pm$ 0.03	2.91 $\pm$ 0.33
weight+TTT	0.41 $\pm$ 0.14 ⁽³⁾	15.72 $\pm$ 2.10	22.24 $\pm$ 3.17	2.00 $\pm$ 0.05	0.14 $\pm$ 0.02	8.03 $\pm$ 0.30
RESfM (sup.)	0.37 $\pm$ 0.12	10.50 $\pm$ 1.45	19.16 $\pm$ 3.42	0.39 $\pm$ 0.72	0.35 $\pm$ 0.04	7.37 $\pm$ 2.63
Rotation error ($^\circ$, mean $\pm$ std)						
remove	4.04 $\pm$ 0.60	6.94 $\pm$ 1.76	12.27 $\pm$ 2.84	14.39 $\pm$ 1.03	16.04 $\pm$ 4.71	13.89 $\pm$ 1.44
remove+TTT	4.22 $\pm$ 0.36	6.90 $\pm$ 1.85	16.40 $\pm$ 1.09	13.23 $\pm$ 0.78	13.19 $\pm$ 6.05	14.53 $\pm$ 2.01
hybrid	3.32 $\pm$ 0.76 ⁽³⁾	8.32 $\pm$ 1.74 ⁽³⁾	28.75 ⁽¹⁾	14.42 $\pm$ 1.99 ⁽³⁾	8.83 $\pm$ 1.46 ⁽³⁾	12.54 ⁽¹⁾
hybrid+TTT	3.24 $\pm$ 0.30 ⁽³⁾	10.85 $\pm$ 0.90 ⁽³⁾	32.53 ⁽¹⁾	14.64 $\pm$ 1.72 ⁽³⁾	8.23 $\pm$ 1.50 ⁽³⁾	13.76 ⁽¹⁾
madweight	3.70 $\pm$ 0.30	7.74 $\pm$ 0.94	27.65 $\pm$ 8.01	12.80 $\pm$ 2.57	9.01 $\pm$ 2.37	14.44 $\pm$ 0.51
weight	3.28 $\pm$ 0.49	8.24 $\pm$ 1.97	27.95 $\pm$ 3.55	7.61 $\pm$ 0.31	9.17 $\pm$ 2.22	13.24 $\pm$ 0.94
weight+TTT	2.98 $\pm$ 0.68 ⁽³⁾	10.09 $\pm$ 1.30	30.60 $\pm$ 3.78	7.95 $\pm$ 0.73	8.30 $\pm$ 1.21	11.60 $\pm$ 1.19
RESfM (sup.)	1.89 $\pm$ 0.46	9.62 $\pm$ 0.63	27.30 $\pm$ 5.06	1.56 $\pm$ 2.54	31.66 $\pm$ 4.19	12.62 $\pm$ 0.82
Registered cameras N_r (% , mean $\pm$ std)						
remove	71.9 $\pm$ 0.9	58.0 $\pm$ 1.2	28.8 $\pm$ 0.9	97.4 $\pm$ 1.7	87.8 $\pm$ 2.0	78.3 $\pm$ 1.4
remove+TTT	70.7 $\pm$ 1.8	57.5 $\pm$ 2.3	32.0 $\pm$ 1.1	95.6 $\pm$ 1.7	89.7 $\pm$ 2.0	78.3 $\pm$ 1.2
hybrid	76.8 $\pm$ 1.0 ⁽³⁾	81.9 $\pm$ 0.4 ⁽³⁾	41.2 ⁽¹⁾	71.6 $\pm$ 4.7 ⁽³⁾	91.9 $\pm$ 1.9 ⁽³⁾	87.5 ⁽¹⁾
hybrid+TTT	77.5 $\pm$ 0.6 ⁽³⁾	81.3 $\pm$ 0.9 ⁽³⁾	43.2 ⁽¹⁾	72.2 $\pm$ 6.4 ⁽³⁾	94.1 $\pm$ 0.7 ⁽³⁾	88.6 ⁽¹⁾
madweight	81.4 $\pm$ 1.1	69.8 $\pm$ 2.4	42.6 $\pm$ 4.9	91.5 $\pm$ 4.1	92.2 $\pm$ 1.7	85.7 $\pm$ 1.9
weight	78.5 $\pm$ 1.4	79.0 $\pm$ 2.9	40.2 $\pm$ 2.3	97.8 $\pm$ 1.5	93.5 $\pm$ 3.3	89.9 $\pm$ 1.2
weight+TTT	77.5 $\pm$ 0.8 ⁽³⁾	81.2 $\pm$ 1.3	41.2 $\pm$ 3.1	98.5 $\pm$ 1.5	94.4 $\pm$ 0.9	90.0 $\pm$ 0.4
RESfM (sup.)	84.6 $\pm$ 1.2	77.3 $\pm$ 1.6	34.3 $\pm$ 3.6	99.3 $\pm$ 1.0	91.7 $\pm$ 1.8	89.3 $\pm$ 1.5

Table 16. Extended removal-statistic ablation, five-seed bands (extends main-paper Table 6; head reweighting fixed). Same conventions as Table 15.

Statistic	MD	1DSfM	1DSfM-hard	Strecha	BMVS	Olsson
Translation error (mean $\pm$ std over training seeds)						
MAD (madweight)	0.51 $\pm$ 0.08	11.04 $\pm$ 1.68	18.92 $\pm$ 4.05	2.65 $\pm$ 0.25	0.14 $\pm$ 0.02	3.01 $\pm$ 0.34
Huber	0.60 $\pm$ 0.03 ⁽³⁾	11.49 $\pm$ 1.63 ⁽³⁾	16.85 ⁽¹⁾	2.73 $\pm$ 0.11 ⁽³⁾	0.14 $\pm$ 0.03 ⁽³⁾	3.04 ⁽¹⁾
STD	0.52 $\pm$ 0.07 ⁽³⁾	15.23 $\pm$ 1.21 ⁽³⁾	23.22 ⁽¹⁾	2.36 $\pm$ 0.24 ⁽³⁾	0.25 $\pm$ 0.09 ⁽³⁾	2.93 ⁽¹⁾
Rotation error ($^\circ$, mean $\pm$ std)						
MAD (madweight)	3.70 $\pm$ 0.30	7.74 $\pm$ 0.94	27.65 $\pm$ 8.01	12.80 $\pm$ 2.57	9.01 $\pm$ 2.37	14.44 $\pm$ 0.51
Huber	4.44 $\pm$ 0.23 ⁽³⁾	7.60 $\pm$ 0.36 ⁽³⁾	24.48 ⁽¹⁾	13.66 $\pm$ 1.94 ⁽³⁾	8.64 $\pm$ 2.07 ⁽³⁾	15.12 ⁽¹⁾
STD	3.83 $\pm$ 1.15 ⁽³⁾	8.98 $\pm$ 0.20 ⁽³⁾	27.34 ⁽¹⁾	9.92 $\pm$ 2.36 ⁽³⁾	17.23 $\pm$ 6.87 ⁽³⁾	13.44 ⁽¹⁾
Registered cameras N_r (% , mean $\pm$ std)						
MAD (madweight)	81.4 $\pm$ 1.1	69.8 $\pm$ 2.4	42.6 $\pm$ 4.9	91.5 $\pm$ 4.1	92.2 $\pm$ 1.7	85.7 $\pm$ 1.9
Huber	80.7 $\pm$ 2.6 ⁽³⁾	70.9 $\pm$ 0.8 ⁽³⁾	38.3 ⁽¹⁾	92.6 $\pm$ 1.9 ⁽³⁾	90.7 $\pm$ 7.5 ⁽³⁾	85.5 ⁽¹⁾
STD	78.9 $\pm$ 0.5 ⁽³⁾	75.1 $\pm$ 1.5 ⁽³⁾	41.8 ⁽¹⁾	91.4 $\pm$ 3.9 ⁽³⁾	91.7 $\pm$ 2.4 ⁽³⁾	87.9 ⁽¹⁾

Table 17. Extended per-scene selection (extends main-paper Table 7; mean translation, then rotation, then N_r ; seed 20). The oracle (italics) selects by ground-truth translation error over five mechanisms; its rotation and N_r are those of the selected arms.

Dataset	weight	madwt	contam	oracle (GT)
Translation error (means)				
MegaDepth	0.423	0.422	0.422	<i>0.116</i>
1DSfM	12.79	12.39	12.39	<i>5.78</i>
Strecha	2.00	2.50	2.00	<i>1.98</i>
BlendedMVS	0.100	0.146	0.100	<i>0.083</i>
Rotation error ($^\circ$, means)				
MegaDepth	3.04	3.49	3.49	<i>1.22</i>
1DSfM	8.39	8.31	8.31	<i>5.50</i>
Strecha	7.48	14.62	7.48	<i>7.41</i>
BlendedMVS	6.65	9.60	6.65	<i>5.28</i>
Registered cameras N_r				
MegaDepth	79%	80%	80%	<i>75%</i>
1DSfM	77%	74%	74%	<i>66%</i>
Strecha	96%	91%	96%	<i>98%</i>
BlendedMVS	88%	95%	88%	<i>90%</i>